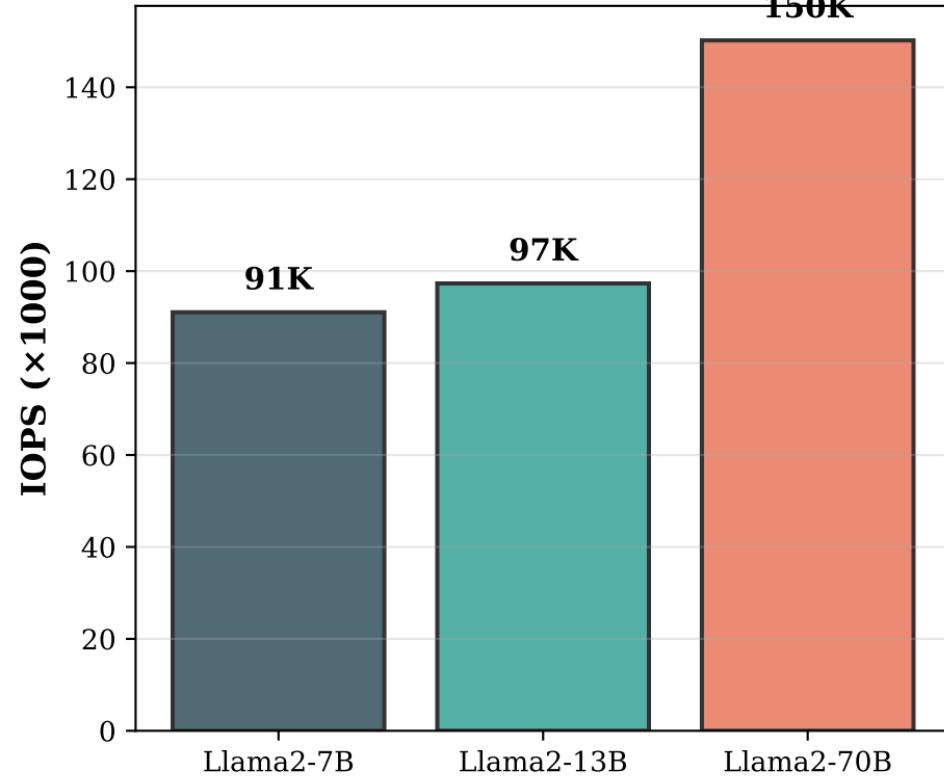


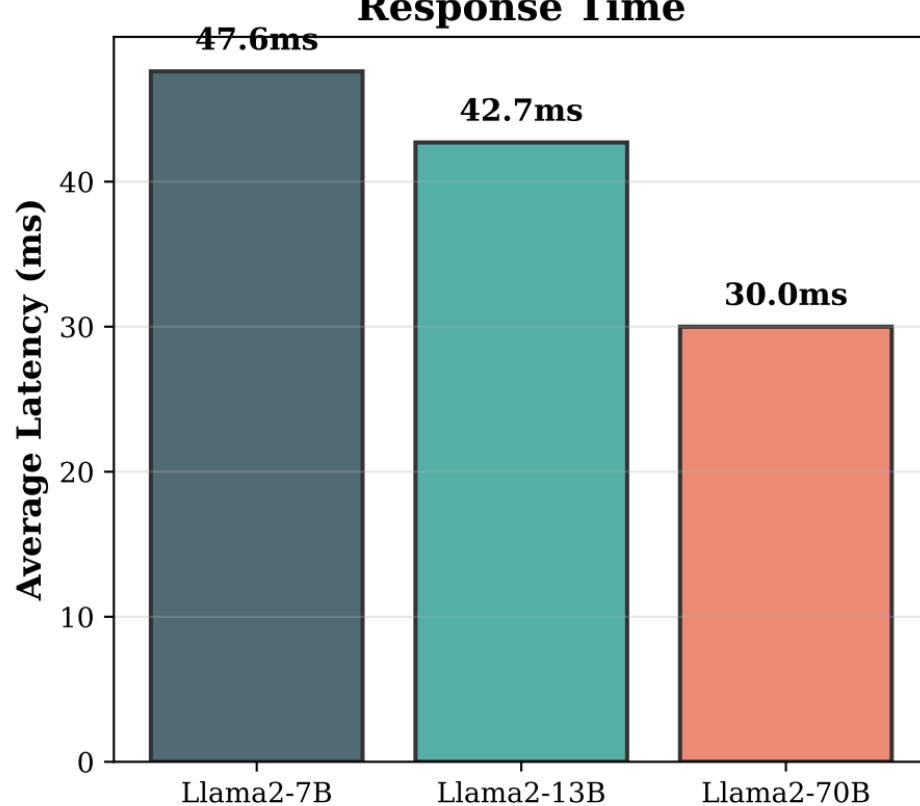
# Baseline Performance Comparison (10K tokens)

Larger Models → Better Parallelism

Throughput



Response Time



Reliability

